machine learning and the market

Justin Fischer
GA DSI11-DEN

problem statement

- Predict short term price changes for publicly traded stocks
  - Specifically the 30 companies in the DJIA
  - Just using price and volume data combined with twitter sentiment

efficient markets hypothesis

- The price of a stock reflects all known information
- 3 Forms:
  - Weak
    - No returns earned based on historical prices; No patterns
  - Semi-Strong
    - Share prices adjust quickly and unbiased, no returns can be earned by trading on new information
  - Strong
    - Reflect all information; No one can earn excess returns
    - Impossible if insider trading laws exist

data

- Downloaded tick data for each company in the Dow Jones Industrial Average
- 9/3/19 - 12/31/19
- Removed weekends and non-trading hours
- Pulled tweets for the same time period using a query that is ‘$’+ticker

making the twitter scraper work for me

- GetOldTweets3 kept timing out and giving too many requests errors
- Had to add a few lines to the library to make it sleep and retry when errors were detected
- Used the time.sleep() function and a ratelimit library

data processing

- Calculated a sentiment score for each tweet
- Resampled each dataset by minute and second
- Calculated a weighted mean for the sentiment score
- Calculated percent changes for the high and mean price in each period
    - 1, 2, 3, 4, 5, 10, 15, 30, 60

quick data summary

	average	low	high
tweet count	11,149.35	1,321 (trv)	81,967 (aapl)
trade count	255,969.5	45,818 (trv)	779,963 (aapl)

Average sentiment - .127

target summary

	periods	down	flat	up
seconds	2,012,484	2.5%	95%	2.5%
minutes	33,625	38.7%	22.3%	39%

goal

- The actual price does not matter here at all
- Predicting a short term movement in a direction is enough to build an effective trading model off of
    - Think *Flash Boys* by Michael Lewis instead of Warren Buffett

- Created a target variable based on the percentage change from the previous period

model

- Recurrent Neural Network on time series data
  - Used because it excels at sequential data
- Parameters:
  - Length - 5
  - Batch - 512
  - Hidden Layers:
    - GRU - 32
    - GRU - 16
    - Dense - 8
    - Dense - 4
  - 100 epochs

Sigh....unsaved results

- After testing it, for some reason the history object didn't save.
- Accuracy numbers may not be the most accurate

seconds

- From quick analysis of the results that did save, since the target was so imbalanced, it just picked ‘flat’ for each tick.
- A ‘flat’ prediction means we don’t do anything since you can’t make money on a stock that doesn’t move
- So, even an accuracy score of close of 95% means nothing since we just predict it to stay the same

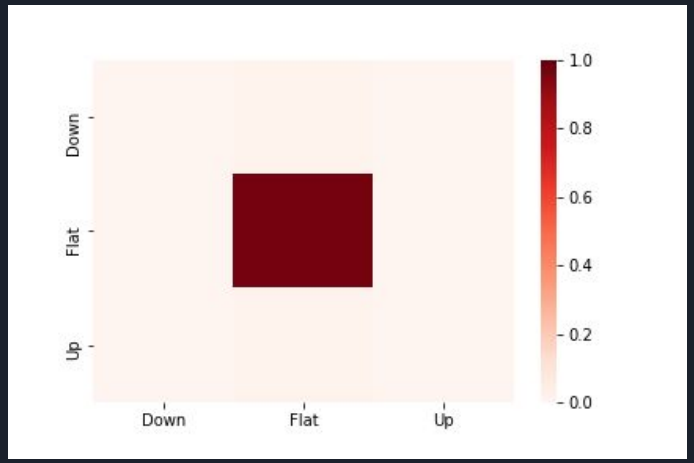

minutes

- We have stopped picking flat periods
- Still are not perfect or close enough since each corner is very similar

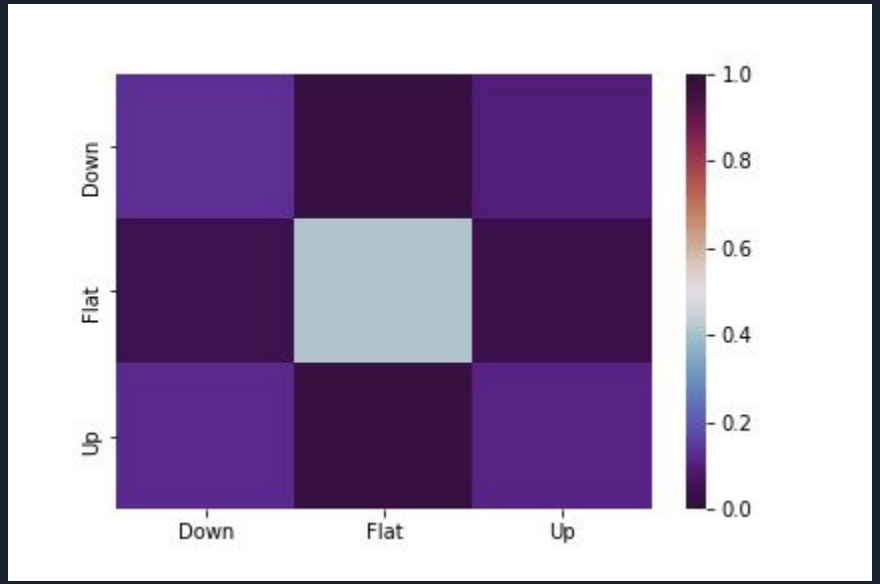

accuracy

- Most companies have a 60-70% accuracy score on the test data
- A few still had the potential to continue to increase with more iterations
- Others were very unpredictable and bounced around

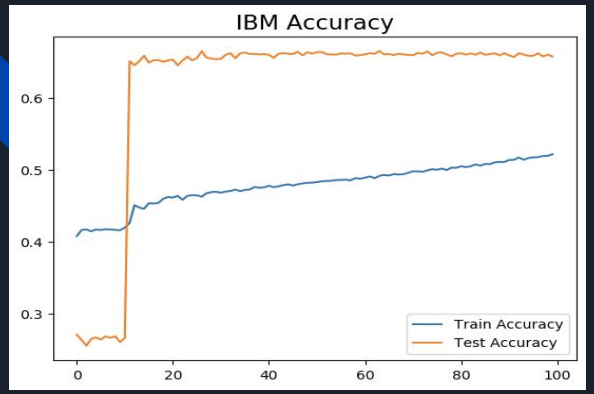
IBM Accuracy
Train Accuracy
Test Accuracy

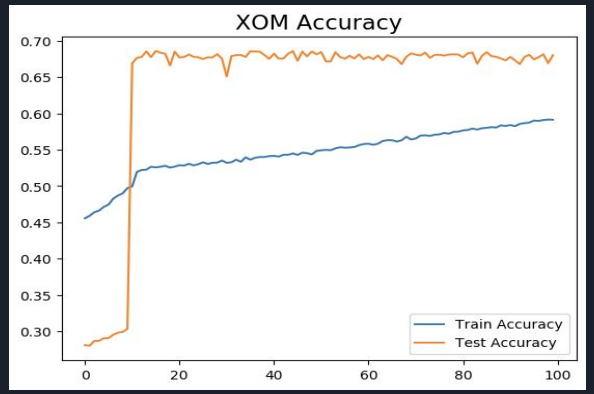
XOM Accuracy
Train Accuracy
Test Accuracy

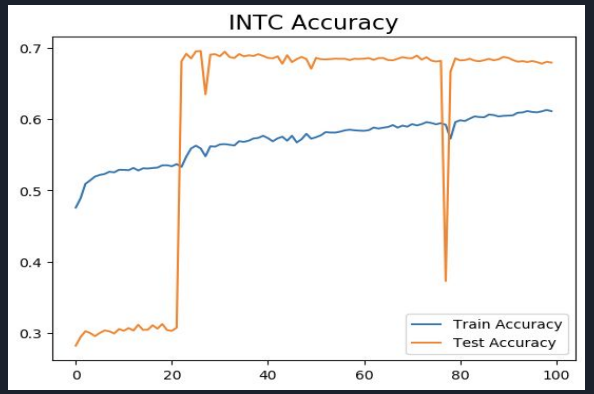
INTC Accuracy
Train Accuracy
Test Accuracy

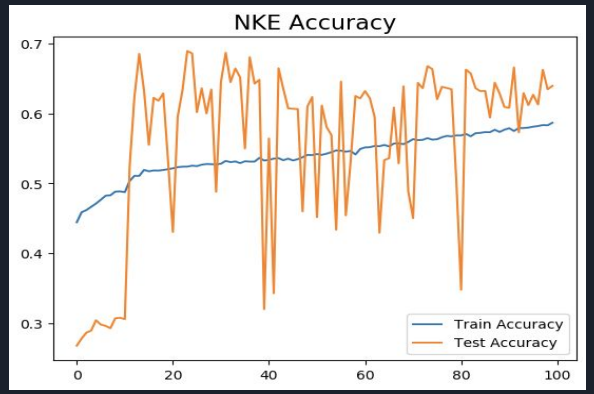
NKE Accuracy
Train Accuracy
Test Accuracy

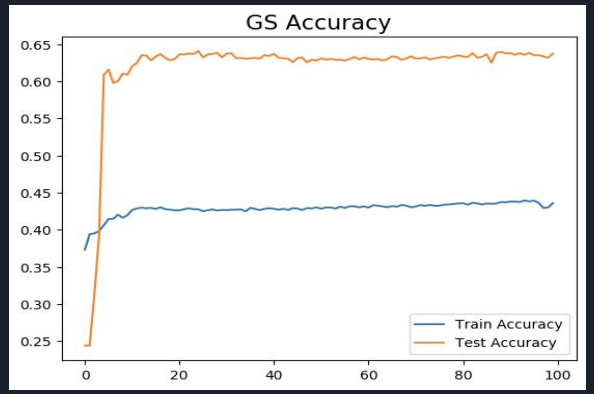
GS Accuracy
Train Accuracy
Test Accuracy

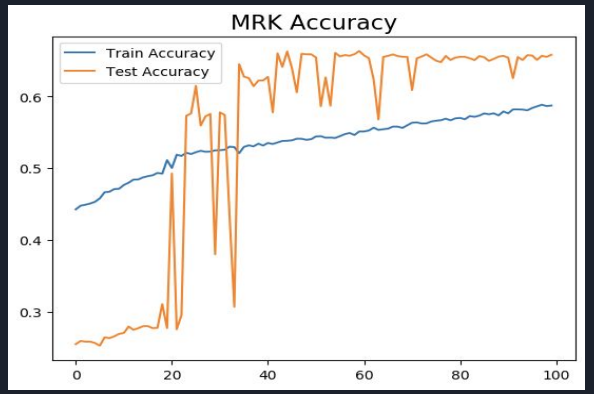
MRK Accuracy
Train Accuracy
Test Accuracy

loss

- The shapes are really strange
- They seem to have a slight curve still, a number of the models could still get better
- Unclear how much better we can make the models with the current parameters

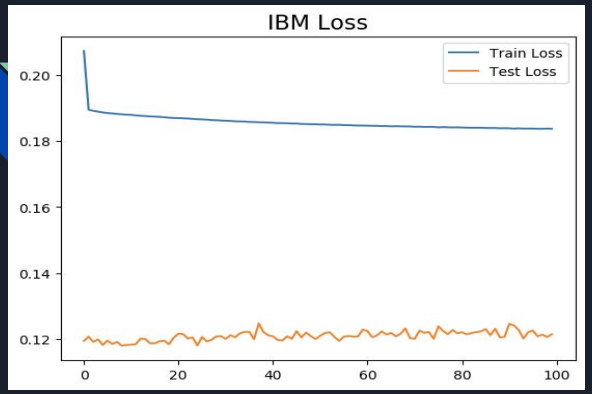
IBM Loss
Train Loss
Test Loss
0.20
0.18
0.16
0.14
0.12
0
20
40
60
80
100

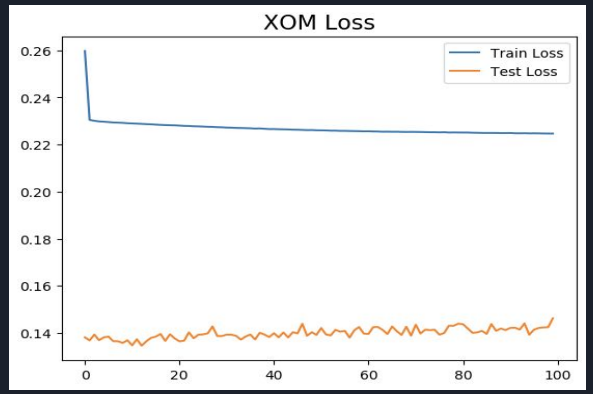
XOM Loss
Train Loss
Test Loss
0.26
0.24
0.22
0.20
0.18
0.16
0.14
0
20
40
60
80
100

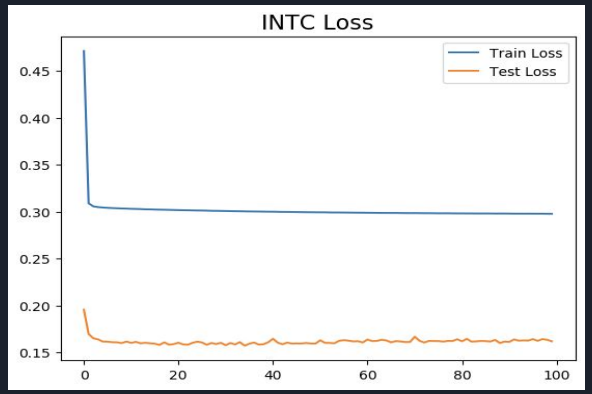
INTC Loss
Train Loss
Test Loss
0.45
0.40
0.35
0.30
0.25
0.20
0.15
0
20
40
60
80
100

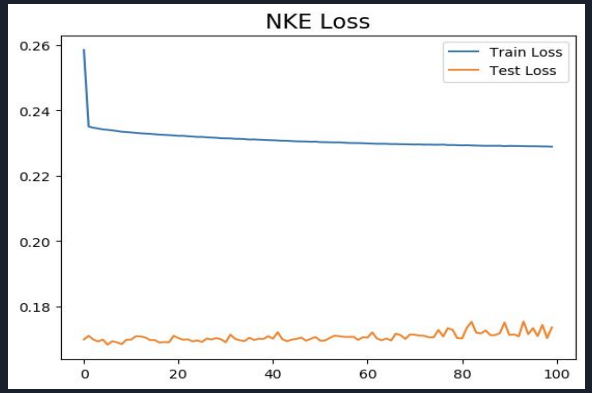
NKE Loss
Train Loss
Test Loss
0.26
0.24
0.22
0.20
0.18
0
20
40
60
80
100

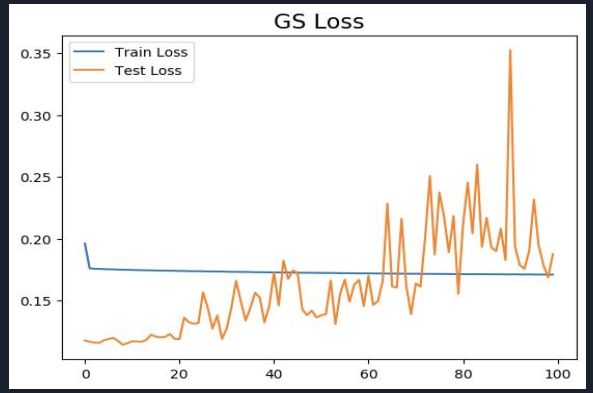
GS Loss
Train Loss
Test Loss
0.35
0.30
0.25
0.20
0.15
0
20
40
60
80
100

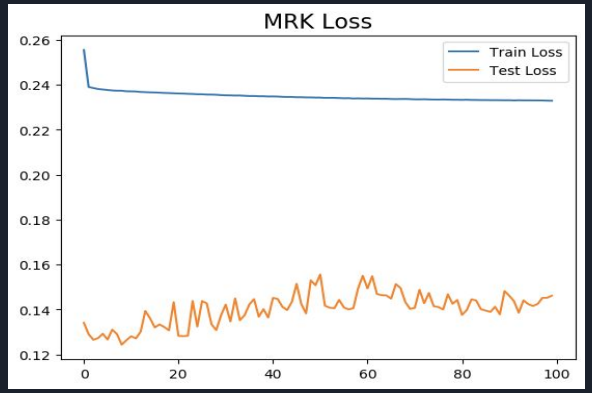
MRK Loss
Train Loss
Test Loss
0.26
0.24
0.22
0.20
0.18
0.16
0.14
0.12
0
20
40
60
80
100

conclusion

- Most of the correct predictions are for no movement
  - This makes it really hard to act upon
- Without the ability to really dig into the results, it is hard to have a firm conclusion
- It seems like it isn't accurate enough to run with money at this time
- Another training session with a slightly modified dataset could be the key

change prediction

- Predict multiple periods, instead of one
  - If prediction is flat, check next period to remake that prediction

next steps

- Different resamples/intervals
  - Combination of a few
- Lesser known and smaller stocks
- Different type of security
  - Commodities
- Use other data
  - Foreign exchange rates
  - Different volume data
  - Bid/Ask spreads and depth
- Better language processing for news and tweets
- Add weights to twitter accounts for those who either:
  - Run a company
  - Famous/respected analysts/talking heads
- Cross-evaluate companies instead of just looking at its own prices
- Try an LSTM model